

2. The SVD.

$$x_1 = \begin{bmatrix} 3 \\ 0 \end{bmatrix}, x_2 = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, x_3 = \begin{bmatrix} 0 \\ \sqrt{6} \end{bmatrix}$$

$$X = [x_1 \ x_2 \ x_3] \in \mathbb{R}^{2 \times 3}$$

$$(a.) \frac{-\sqrt{6}}{6} x_1 + \frac{\sqrt{6}}{2} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = x_3$$

$$u_1 = \frac{x_1}{\|x_1\|_2} = \begin{bmatrix} 3 \\ 0 \end{bmatrix} / 3 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\begin{aligned} x_2' &= x_2 - u_1(u_1^T x_2) = \begin{bmatrix} 1 \\ 2 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \end{bmatrix} \left(\begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} \right) \\ &= \begin{bmatrix} 1 \\ 2 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \end{bmatrix} \\ &= \begin{bmatrix} 0 \\ 2 \end{bmatrix} \end{aligned}$$

$$u_2 = \frac{x_2'}{\|x_2'\|_2} = \begin{bmatrix} 0 \\ 2 \end{bmatrix} / 2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\Rightarrow U = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

(b.)

(i) $\in \mathbb{R}^2$

$$V = \text{span}(v)$$

$$P_V = v(v^T v)^{-1} v^T$$

$$\begin{aligned} P_V^T &= v(v^T(v^T)^{-1})^T v^T \\ &= v(v^T(v^T)^{-1})^T v^T \\ &= v(v^T(v^T)^{-1}) v^T \\ &= v(v^T v)^{-1} v^T = P_V \end{aligned}$$

(ii)

$$d_i^2 = \|x_i - P_V x_i\|_2^2 = (x_i - P_V x_i)^T (x_i - P_V x_i)$$

$$= x_i^T x_i - x_i^T P_V^T x_i - x_i^T P_V x_i + x_i^T P_V^T P_V x_i$$

$$= \|x_i\|_2^2 - 2 x_i^T P_V x_i + x_i^T P_V^2 x_i$$

$$= \|x_i\|_2^2 - 2 x_i^T P_V x_i + x_i^T P_V x_i$$

$$= \|x_i\|_2^2 - x_i^T P_V x_i$$

(iii)

$$d_i^2 = \|x_i\|_2^2 - x_i^T p_i x_i$$

$$v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}, \quad x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \end{bmatrix}$$

$$x_i^T p_i x_i = x_i^T v (v^T v)^{-1} v^T x_i$$

$$= x_i^T \frac{v v^T}{(v^T v)} x_i$$

$$= \frac{(v^T x_i)^2}{v^T v}$$

$$= \frac{(v_1 x_{i1} + v_2 x_{i2})^2}{v_1^2 + v_2^2}$$

$$\Rightarrow d_i^2 = x_i^T x_i - \frac{(v_1 x_{i1} + v_2 x_{i2})^2}{v_1^2 + v_2^2}$$

(C)

$$X = \begin{bmatrix} 3 & 1 & 0 \\ 0 & 2 & \sqrt{6} \end{bmatrix}$$

$$\arg \min_v \sum_{i=1}^3 d_i^2 = \arg \max_{v: v^T v = 1} \sum_{i=1}^3 \frac{(v_1 x_{i1} + v_2 x_{i2})^2}{v_1^2 + v_2^2}$$

$$= \arg \max_{v: v^T v = 1} \sum_{i=1}^3 (v_1 x_{i1} + v_2 x_{i2})^2$$

$$\begin{aligned} \sum_{i=1}^3 (v_1 x_{i1} + v_2 x_{i2})^2 &= (3v_1)^2 + (v_1 + 2v_2)^2 + (\sqrt{6}v_2)^2 \\ &= 9v_1^2 + v_1^2 + 4v_1v_2 + 4v_2^2 + 6v_2^2 \\ &= 10v_1^2 + 4v_1v_2 + 10v_2^2 \\ &= 10(v_1^2 + v_2^2) + 4v_1v_2 \\ &= 10 + 4v_1v_2 \end{aligned}$$

$$\mathcal{L} = 10 + 4v_1v_2 - \lambda(v_1^2 + v_2^2 - 1)$$

$$\frac{\partial \mathcal{L}}{\partial v_1} = 4v_2 - 2\lambda v_1 = 0$$

$$\Rightarrow v_1 = v_2 \Rightarrow 2v_1^2 = 1 \Rightarrow v_1 = v_2 = \frac{1}{\sqrt{2}} \Rightarrow v = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$$

$$\frac{\partial \mathcal{L}}{\partial v_2} = 4v_1 - 2\lambda v_2 = 0$$

$$u_1 = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}, \quad u_2 = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} \end{bmatrix} \Rightarrow U = \begin{bmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} \end{bmatrix}$$

Include part of (iv)

(iv) $v' = \begin{bmatrix} -\frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$ also minimize

d_i^2 . So v is not unique.

But p is unique as

$$v v^T = v' v'^T.$$

$$\sigma_1 = \|X^T U_1\|_2$$

$$= \left\| \begin{bmatrix} 3 & 0 \\ 1 & 2 \\ 0 & \sqrt{6} \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} \end{bmatrix} \right\|_2 = \left\| \begin{bmatrix} \frac{3}{\sqrt{2}} \\ \frac{3}{\sqrt{2}} \\ \frac{\sqrt{6}}{\sqrt{2}} \end{bmatrix} \right\|_2 = \sqrt{\frac{9}{2} + \frac{9}{2} + \frac{6}{2}} = \sqrt{12} = 2\sqrt{3}$$

$$\sigma_2 = \|X^T U_2\|_2$$

$$= \left\| \begin{bmatrix} 3 & 0 \\ 1 & 2 \\ 0 & \sqrt{6} \end{bmatrix} \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix} \right\|_2 = \left\| \begin{bmatrix} \frac{3}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} \\ -\frac{\sqrt{6}}{\sqrt{2}} \end{bmatrix} \right\|_2 = \sqrt{\frac{9}{2} + \frac{1}{2} + \frac{6}{2}} = \sqrt{8} = 2\sqrt{2}$$

$$\Rightarrow X \in \mathbb{R}^{2 \times 3}$$

$$\Rightarrow \Sigma = \begin{bmatrix} 2\sqrt{3} & 0 & 0 \\ 0 & 2\sqrt{2} & 0 \end{bmatrix}$$

$$3. X \in \mathbb{R}^{n \times d} \quad U = [u_1 \ u_2 \ \dots \ u_r] \quad V = [v_1 \ v_2 \ \dots \ v_r] \quad X = \underbrace{U}_{n \times r} \underbrace{\Sigma}_{r \times r} \underbrace{V^T}_{r \times d}$$

$$(a) \quad X = \sum_{i=1}^r \sigma_i u_i v_i^T$$

$$(b) \quad X_k = \sum_{i=1}^k \sigma_i u_i v_i^T$$

$$4. \quad X = \begin{bmatrix} 5 & 0 \\ 0 & 1 \end{bmatrix}$$

X is a diagonal matrix

$$\Rightarrow U = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad \Sigma = \begin{bmatrix} 5 & 0 \\ 0 & 1 \end{bmatrix}, \quad V = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$X = U \Sigma V = I \Sigma I^T = \Sigma$$

$$5. \quad X = \begin{bmatrix} -3 & 0 \\ 0 & -1 \end{bmatrix}$$

X is a diagonal matrix

$$\Rightarrow U = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}, \quad \Sigma = \begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix}, \quad V = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$X = U \Sigma V = -1 I \Sigma I^T = -\Sigma = \begin{bmatrix} -3 & 0 \\ 0 & -1 \end{bmatrix}$$

cmsc253_hw4

November 5, 2023

```
[ ]: import numpy as np
import numpy.linalg as la # matrix rank, inverse import scipy.io
import matplotlib.pyplot as plt # plots
import scipy.io as sio
import sys
from matplotlib import colors
```

Q1 Gram-Schmidt. (a)

```
[ ]: def gram_schmidt(X):
    # X is an n-by-p matrix.
    # Returns U an orthonormal matrix.
    # eps is a threshold value to identify if a vector # is nearly a zero
    ↪vector.

    eps = 1e-12

    n, p = X.shape
    U = np.zeros((n, 0))
    for j in range(p):
        # Get the j-th column of matrix X
        v = X[:, j]
        # Write your own code here: Perform the
        # orthogonalization by subtracting the projections on
        # all columns of U. And then check whether the vector
        # you get is nearly a zero vector.
        v = v - U@U.T@v
        v = np.reshape(v, (-1, 1))
        if np.linalg.norm(v) > eps:
            # Normalize vector v and append it to U
            U = np.hstack((U, v / la.norm(v)))
    return U
X = np.array([[1, 1/2, 1/3, 1/4, 1/5, 1/6, 1/7],
              [1/2, 1/3, 1/4, 1/5, 1/6, 1/7, 1/8],
              [1/3, 1/4, 1/5, 1/6, 1/7, 1/8, 1/9],
              [1/4, 1/5, 1/6, 1/7, 1/8, 1/9, 1/10],
              [1/5, 1/6, 1/7, 1/8, 1/9, 1/10, 1/11],
              [1/6, 1/7, 1/8, 1/9, 1/10, 1/11, 1/12],
```

```
[1/7, 1/8, 1/9, 1/10, 1/11, 1/12, 1/13]])
gram_schmidt(X)
```

```
[ ]: array([[ 8.13304645e-01, -5.43794290e-01,  1.99095308e-01,
            -5.51132525e-02,  1.19578493e-02, -1.96870730e-03,
            -2.10219152e-04],
           [ 4.06652323e-01,  3.03317262e-01, -6.88560548e-01,
            4.75965265e-01, -1.97392654e-01,  5.41151990e-02,
            3.01288446e-03],
           [ 2.71101548e-01,  3.93949644e-01, -2.07134626e-01,
            -4.90111167e-01,  6.10839674e-01, -3.26893986e-01,
            1.54577469e-02],
           [ 2.03326161e-01,  3.81744394e-01,  1.12389156e-01,
            -4.39598473e-01, -2.54179063e-01,  6.41045574e-01,
            -2.08187091e-01],
           [ 1.62660929e-01,  3.51412668e-01,  2.91518455e-01,
            -1.12276608e-01, -4.99206689e-01, -2.02224426e-01,
            6.15118567e-01],
           [ 1.35550774e-01,  3.20235052e-01,  3.89191824e-01,
            2.30886766e-01, -1.50594721e-01, -5.41817789e-01,
            -7.06123659e-01],
           [ 1.16186378e-01,  2.92095792e-01,  4.40744903e-01,
            5.20623748e-01,  5.01273338e-01,  3.80535581e-01,
            2.81830797e-01]])
```

(b)

```
[ ]: def hilbert_matrix(n):
      X = np.array([[1.0 / (i + j - 1) for i in \
                      range(1, n + 1)] for j in range(1, n + 1)])
      return X
n = 7
X = hilbert_matrix(n)
U = gram_schmidt(X)
L1 = np.identity((U.T@U).shape[0]) - U.T@U
print(L1)
err = la.norm(L1, ord=1)
print(f'L1 matrix norm is {err}')
```

```
[[ 0.00000000e+00  6.73072709e-16 -2.36616282e-15  2.80331314e-14
   -1.92849903e-12  1.25247715e-10 -9.71088800e-09]
 [ 6.73072709e-16  0.00000000e+00 -2.76723089e-14  3.71008779e-13
  -2.63988831e-12 -4.05814549e-11  5.49175827e-09]
 [-2.36616282e-15 -2.76723089e-14  1.11022302e-16  8.52865001e-12
  -2.31620473e-10  5.20528393e-09 -1.33284293e-07]
 [ 2.80331314e-14  3.71008779e-13  8.52865001e-12 -2.22044605e-16
  -9.04092928e-09  4.29613409e-07 -1.80981921e-05]
 [-1.92849903e-12 -2.63988831e-12 -2.31620473e-10 -9.04092928e-09
```

```

0.00000000e+00 2.66371580e-05 -2.30322707e-03]
[ 1.25247715e-10 -4.05814549e-11 5.20528393e-09 4.29613409e-07
 2.66371580e-05 0.00000000e+00 -2.27098006e-01]
[-9.71088800e-09 5.49175827e-09 -1.33284293e-07 -1.80981921e-05
-2.30322707e-03 -2.27098006e-01 1.11022302e-16]]
L1 matrix norm is 0.22941947980869065

```

(c)

```

[ ]: def modified_gram_schmidt(X):
    # Define a threshold value to identify if a vector # is nearly a zero
    ↪vector.
    eps = 1e-12
    n, p = X.shape
    U = np.zeros((n, 0))
    for j in range(p):
        # Get the j-th column of matrix X
        v = X[:, j]
        for i in range(j):
            # Compute and subtract the projection of
            # vector v onto the i-th column of U
            v = v - np.dot(U[:, i], v) * U[:, i]
        v = np.reshape(v, (-1, 1))
        # Check whether the vector we get is nearly # a zero vector
        if np.linalg.norm(v) > eps:
            # Normalize vector v and append it to U
            U = np.hstack((U, v / la.norm(v)))
    return U

n = 7
X = hilbert_matrix(n)
U_modified = modified_gram_schmidt(X)
L1_modified = np.identity((U_modified.T@U_modified).shape[0]) - \
    U_modified.T@U_modified

L1_modified
print(L1_modified)
err_modified = la.norm(L1_modified, ord=1)
print(f'Modified L1 matrix norm is {err_modified}')

```

```

[[ 0.00000000e+00 1.54737334e-15 -2.16146545e-14 3.62931907e-13
 -7.97459321e-12 3.16995839e-10 -1.90386089e-08]
 [ 1.54737334e-15 0.00000000e+00 1.02695630e-15 -2.63955524e-14
 9.00696184e-13 -2.74882339e-11 1.03917761e-09]
 [-2.16146545e-14 1.02695630e-15 -2.22044605e-16 3.05311332e-15
 -2.71893619e-13 1.51737511e-11 -6.93317848e-10]
 [ 3.62931907e-13 -2.63955524e-14 3.05311332e-15 0.00000000e+00
 -1.66533454e-15 -2.34368080e-13 4.83647289e-11]
 [-7.97459321e-12 9.00696184e-13 -2.71893619e-13 -1.66533454e-15

```



```

0.00000000e+00  2.01227923e-14 -2.16494878e-12]
[ 3.16995839e-10 -2.74882339e-11  1.51737511e-11 -2.34368080e-13
 2.01227923e-14  2.22044605e-16 -1.46549439e-14]
[-1.90386089e-08  1.03917761e-09 -6.93317848e-10  4.83647289e-11
 -2.16494878e-12 -1.46549439e-14 -2.22044605e-16]]
Modified L1 matrix norm is 2.0821648873819987e-08

```

In the traditional approach to Gram-Schmidt, each vector in the sequence is made orthogonal to all previously processed vectors by subtracting their projections. The modified Gram-Schmidt re-orthogonalizes each vector against all previously orthogonalized vectors, one at a time. This step-wise approach reduces the accumulation of errors and maintains better orthogonality. Therefore, the modified Gram-Schmidt process is generally superior to the original process.

Q6 (a)

```

[ ]: d = np.load('face_emotion_data.npz')
X = d['X']
y = d['y']
n, p = np.shape(X)
# error rate for truncated SVD
error_SVD = np.zeros((8, 7))
# SVD parameters to test
k_vals = np.arange(9) + 1
param_err_SVD = np.zeros(len(k_vals))

subset_count = 8

subsets = np.array_split(np.arange(n), 8)

for i in range(subset_count):
    X_hold_out = X[subsets[i]]
    y_hold_out = y[subsets[i]]

    for idx, j in enumerate([j for j in range(subset_count) if j != i]):
        X_train = np.concatenate([X[subsets[p]] for p in range(8) if p != i
                                   and p != j])
        y_train = np.concatenate([y[subsets[p]] for p in range(8) if p != i
                                   and p != j])

        U, Sigma, VT = la.svd(X_train, full_matrices=False)

        X_test = X[subsets[j]]
        y_test = y[subsets[j]]

        w_hat_lst = list()

        for k_idx, k in enumerate(k_vals):
            Sigma_plus = np.diag(1 / Sigma[:k])

```



```

        VT_k = VT[:k, :]
        U_k = U[:, :k]

        w_hat = VT_k.T@Sigma_plus@U_k.T@y_train
        w_hat_lst.append(w_hat)
        y_hat = X_test@w_hat

        param_err_SVD[k_idx] += np.sum(np.sign(y_hat) != y_test) / 16

    # print(param_err_SVD)
    k_best = np.argmin(param_err_SVD)
    w_hat_best = w_hat_lst[k_best]

    y_hat_hold_out = X_hold_out@w_hat_best
    error_SVD[i][idx] = np.sum(np.sign(y_hat_hold_out) != y_hold_out) / 16
print(f"Error estimate: {np.mean(error_SVD)}")

```

Error estimate: 0.052455357142857144

(b)

```

[ ]: d = np.load('face_emotion_data.npz')
X = d['X']
y = d['y']
n, p = np.shape(X)
# error rate for regularized least squares
error_RLS = np.zeros((8, 7))
# RLS parameters to test
lambda_vals = np.array([0, 0.5, 1, 2, 4, 8, 16])
param_err_RLS = np.zeros(len(lambda_vals))

subset_count = 8

subsets = np.array_split(np.arange(n), 8)

for i in range(subset_count):
    X_hold_out = X[subsets[i]]
    y_hold_out = y[subsets[i]]

    for idx, j in enumerate([j for j in range(subset_count) if j != i]):
        X_train = np.concatenate([X[subsets[p]] for p in range(8) if p != i
                                   and p != j])
        y_train = np.concatenate([y[subsets[p]] for p in range(8) if p != i
                                   and p != j])

        X_test = X[subsets[j]]
        y_test = y[subsets[j]]

```

```

w_hat_lst = list()

for lmd_idx, lmd in enumerate(lambda_vals):
    w_hat = la.inv(X_train.T@X_train + lmd*np.identity(p))@X_train.
    ↪T@y_train
    w_hat_lst.append(w_hat)
    y_hat = X_test@w_hat

    param_err_RLS[lmd_idx] += np.sum(np.sign(y_hat) != y_test) / 16

    # print(param_err_SVD)
    k_best = np.argmin(param_err_RLS)
    w_hat_best = w_hat_lst[k_best]

    y_hat_hold_out = X_hold_out@w_hat_best
    error_RLS[i][idx] = np.sum(np.sign(y_hat_hold_out) != y_hold_out) / 16
print(f"Error estimate: {np.mean(error_RLS)}")

```

Error estimate: 0.0546875

(c)

```

[ ]: X = d['X']
y = d['y']
new_features = X@np.random.rand(9,3)
X_new = np.concatenate((X, new_features), axis = 1)

### (a)' ###
n, p = np.shape(X_new)

# error rate for regularized least squares
error_SVD = np.zeros((8, 7))
# SVD parameters to test
k_vals = np.arange(9) + 1
param_err_SVD = np.zeros(len(k_vals))

subset_count = 8

subsets = np.array_split(np.arange(n), 8)

for i in range(subset_count):
    X_hold_out = X_new[subsets[i]]
    y_hold_out = y[subsets[i]]

    for idx, j in enumerate([j for j in range(subset_count) if j != i]):
        X_train = np.concatenate([X_new[subsets[p]] for p in range(8) if p != i
                                   and p != j])
        y_train = np.concatenate([y[subsets[p]] for p in range(8) if p != i

```

```

                                and p != j])

    U, Sigma, VT = la.svd(X_train, full_matrices=False)

    X_test = X_new[subsets[j]]
    y_test = y[subsets[j]]

    w_hat_lst = list()

    for k_idx, k in enumerate(k_vals):
        Sigma_plus = np.diag(1 / Sigma[:k])
        VT_k = VT[:k, :]
        U_k = U[:, :k]

        w_hat = VT_k.T@Sigma_plus@U_k.T@y_train
        w_hat_lst.append(w_hat)
        y_hat = X_test@w_hat

        param_err_SVD[k_idx] += np.sum(np.sign(y_hat) != y_test) / 16

    # print(param_err_SVD)
    k_best = np.argmin(param_err_SVD)
    w_hat_best = w_hat_lst[k_best]

    y_hat_hold_out = X_hold_out@w_hat_best
    error_SVD[i][idx] = np.sum(np.sign(y_hat_hold_out) != y_hold_out) / 16
print(f"Error estimate: {np.mean(error_SVD)}")

```

Error estimate: 0.049107142857142856

```

[ ]: ### (b)' ###
# error rate for regularized least squares
error_RLS = np.zeros((8, 7))
# RLS parameters to test
lambda_vals = np.array([0, 0.5, 1, 2, 4, 8, 16])
param_err_RLS = np.zeros(len(lambda_vals))

subset_count = 8

subsets = np.array_split(np.arange(n), 8)

for i in range(subset_count):
    X_hold_out = X_new[subsets[i]]
    y_hold_out = y[subsets[i]]

    for idx, j in enumerate([j for j in range(subset_count) if j != i]):
        X_train = np.concatenate([X_new[subsets[p]] for p in range(8) if p != i

```

```

                                and p != j])
y_train = np.concatenate([y[subsets[p]] for p in range(8) if p != i
                                and p != j])

X_test = X_new[subsets[j]]
y_test = y[subsets[j]]

w_hat_lst = list()

for lmd_idx, lmd in enumerate(lambda_vals):
    w_hat = la.inv(X_train.T@X_train + lmd*np.identity(p))@X_train.
    ↪T@y_train
    w_hat_lst.append(w_hat)
    y_hat = X_test@w_hat

    param_err_RLS[lmd_idx] += np.sum(np.sign(y_hat) != y_test) / 16

# print(param_err_SVD)
k_best = np.argmin(param_err_RLS)
w_hat_best = w_hat_lst[k_best]

y_hat_hold_out = X_hold_out@w_hat_best
error_RLS[i][idx] = np.sum(np.sign(y_hat_hold_out) != y_hold_out) / 16
print(f"Error estimate: {np.mean(error_RLS)}")

```

Error estimate: 0.05357142857142857

The average error rate of Truncated SVD solution is 0.052455357142857144. The average error rate of Regularized LS is 0.0546875. With new features, Avg error rate of SVD_new is 0.049107142857142856 and Avg error rate of RLS_new is 0.05357142857142857. It seems that these new features are helpful for classification, but not significantly so. One possible explanation is that since the new features are generated as random linear combinations of the original features, they might not necessarily add any meaningful or informative variance to the model that would improve classification.